

✔ Congratulations! You passed!

Grade received 100% Latest Submission Grade 100% To pass 80% or higher

Go to next item

1. What statements are true about cloudant? (Select all that apply)

1 / 1 point

- ☒ Cloudant is based on ApacheCouchDB

✔ Correct

correct

- ☐ Cloudant is a SQL database
- ☒ Cloudant is a NoSQL database

✔ Correct

Correct

- ☐ Cloudant is a very fast and scalable key-value store

- ☒ Cloudant is meant for storing JSON documents effectively

✔ Correct

Correct

- ☐ BigCouch is a tool to inflate storage on CouchDB

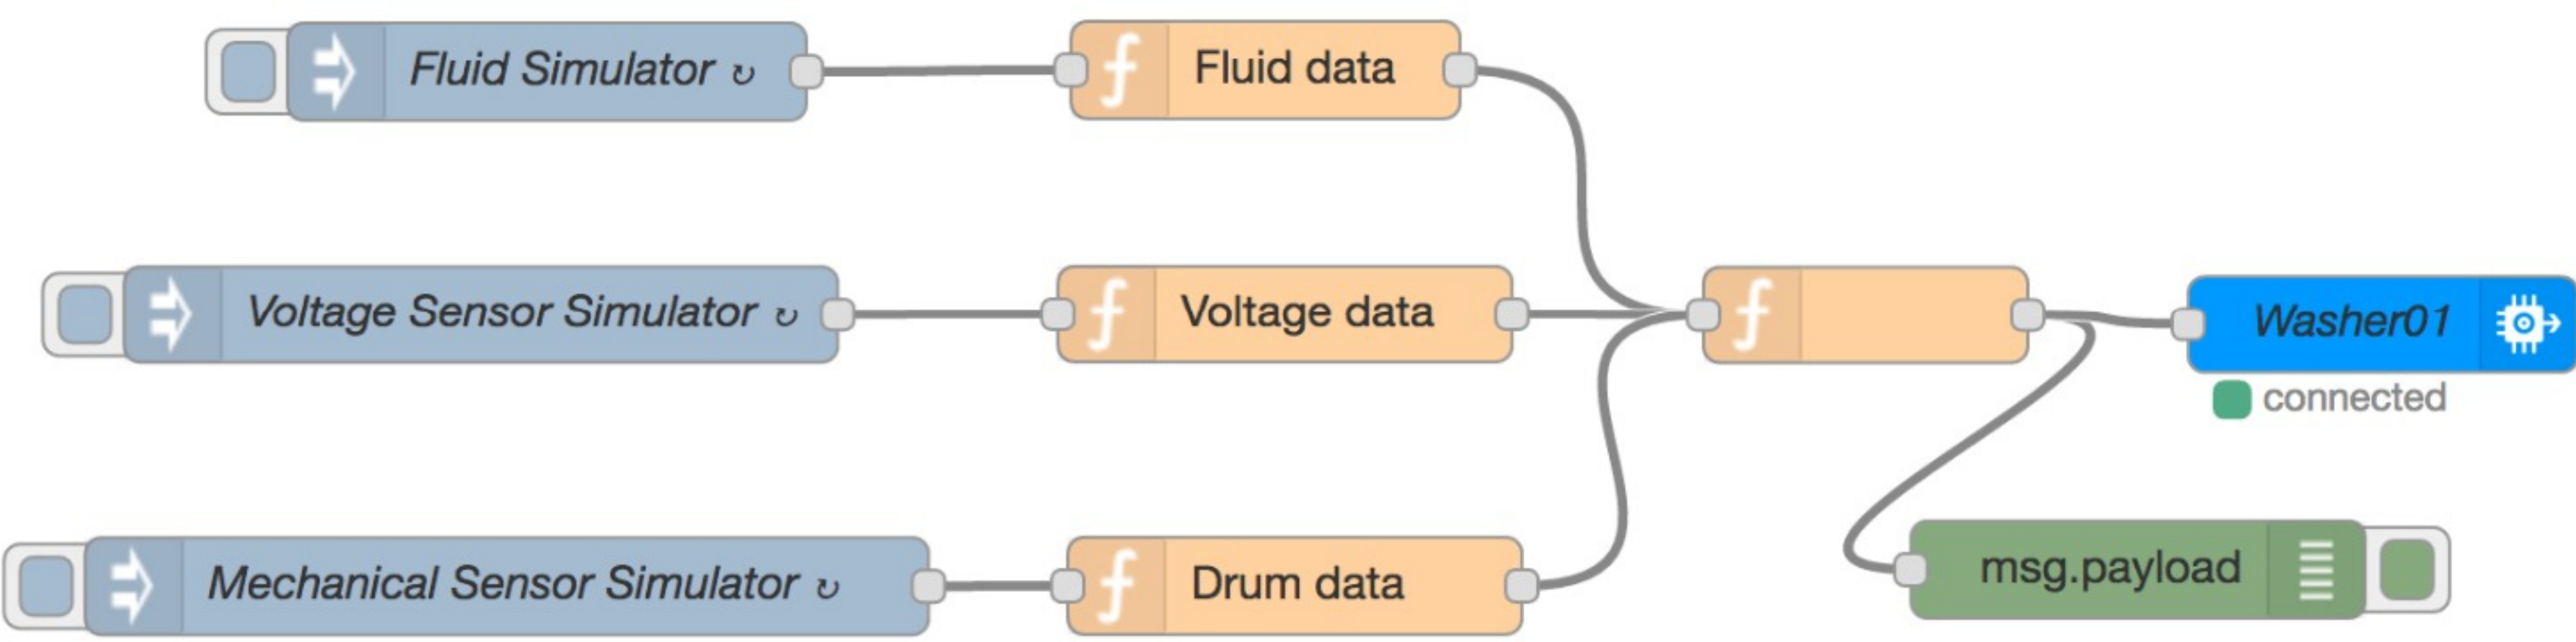
- ☒ BigCouch is a component between the client and a set of CouchDB services used for horizontal scaling

✔ Correct

Correct

2. Please have a look at the following flow:

1 / 1 point



Which nodes are actually simulating sensors of a hypothetical IoT device?

(Please select Fluid Simulator, Voltage Sensor Simulator, Mechanical Sensor Simulator, this is some legacy from a previous version of this course and we can't change the quiz at that point for fairness reasons)

- ☐ Washer01
- ☐ msg.payload
- ☒ Voltage Sensor Simulator

✔ Correct

Correct

- ☐ Voltage Data
- ☒ Fluid Simulator

✔ Correct

Correct

- ☐ Fluid Data
- ☒ Mechanical Sensor Simulator

✔ Correct

Correct

- ☐ Drum Data

3. In the "End-to-End Scenario", where does all the data get stored in?

1 / 1 point

(Please select Cloudant (ApacheCouchDB), this is some legacy from a previous version of this course and we can't change the quiz at that point for fairness reasons)

- ☒ Cloudant (ApacheCouchDB)
- ☐ ApacheSpark
- ☐ Object Storage
- ☐ OpenStack Swift

✔ Correct

Correct

4. How does the Catalyst optimizer work internally?

1 / 1 point

Abbreviations:

AST - Abstract Syntax Tree

LEP - Logical Execution Plan

PEP - Physical Execution Plan

- ☐ A AST is created from an SQL LEP. This AST is transformed (optimised). Then multiple PEPs are created from the optimised LEP. Finnaly, based on cost based statistics an optimal PEP is chosen to be executed.
- ☒ A LEP is created from an SQL AST. This LEP is transformed (optimised). Then multiple PEPs are created from the optimised LEP. Finnaly, based on cost based statistics an optimal PEP is chosen to be executed.
- ☐ A AST is created from an SQL PEP. This AST is transformed (optimised). Then multiple LEPs are created from the optimised PEP. Finnaly, based on cost based statistics an optimal LEP is chosen to be executed.
- ☐ A PEP is created from an SQL AST. This AST is transformed (optimised). Then multiple PEP are created from the optimised LEP. Finnaly, based on cost based statistics an optimal PEP is chosen to be executed.

✔ Correct

Correct

5. What is the advantage of using ApacheSparkSQL over RDDs? (select all that apply)

1 / 1 point

- ☐ ApacheSparkSQL bypasses the RDD interface which has been proven to be very complicated
- ☐ SQL is simpler than RDD but has some performance drawbacks
- ☒ Catalyst and Tungsten are able to optimise the execution, so are more likely to execute more quickly than if you would had implemented something equivalent using the RDD API.

✔ Correct

Correct

- ☒ The API is simpler and doesn't require specific functional programming skills

✔ Correct

Correct